

Ensemble Methods for Corporate Bankruptcy Prediction:

A Comparative Analysis of Random Forest and XGBoost Models
Using Multi-Temporal Financial Data

Advanced Analytics Research
PhD Level Research Report

June 20, 2025

Contents

1	Introduction	1
2	Exploratory Data Analysis	3
3	Modeling Methodology and Results	10
4	Discussion and Conclusions	19

Abstract

This comprehensive study investigates the application of ensemble machine learning methods for corporate bankruptcy prediction using a multi-temporal financial dataset spanning five-year prediction horizons. The research develops and compares four distinct modeling approaches: Random Forest and XGBoost algorithms applied to both comprehensive and selected feature sets derived from 64 financial indicators across 43,405 company observations.

The analysis encompasses rigorous exploratory data analysis, including systematic missing value treatment through K-Nearest Neighbors imputation, correlation-based feature selection using hierarchical clustering, and Variance Inflation Factor validation. Comprehensive hyperparameter optimization employing cross-validation techniques ensures robust model development and fair performance comparison.

Results demonstrate the superior performance of XGBoost with full features, achieving an Area Under the Curve (AUC) of 0.91, significantly outperforming alternative approaches. The study reveals that comprehensive feature sets substantially outperform reduced feature sets, challenging conventional assumptions about feature selection necessity for tree-based ensemble methods. XGBoost consistently demonstrates superior minority class detection capabilities, achieving 36% recall and 89% precision for bankruptcy identification compared to Random Forest's limitations in class imbalance scenarios.

The findings provide definitive guidance for practitioners implementing bankruptcy prediction systems, establishing XGBoost with comprehensive feature sets as the optimal approach for financial distress modeling. This research contributes valuable insights to the growing literature on ensemble methods in financial analytics and offers practical recommendations for financial institutions seeking robust bankruptcy prediction capabilities.

Keywords: Bankruptcy prediction, ensemble methods, XGBoost, Random Forest, feature selection, financial analytics, machine learning, risk assessment

Chapter 1

Introduction

Corporate bankruptcy prediction has emerged as a critical area of research in financial analytics, offering substantial value to stakeholders including investors, creditors, and policymakers. The ability to accurately predict bankruptcy enables informed decision-making, risk assessment, and proactive financial management strategies. This study leverages a comprehensive bankruptcy dataset spanning multiple time horizons to develop robust ensemble models for bankruptcy prediction.

The dataset utilized in this research comprises bankruptcy information for companies across five different time periods (1-year, 2-year, 3-year, 4-year, and 5-year horizons). Each temporal dataset contains 64 financial features that capture various aspects of company performance, including liquidity ratios, profitability metrics, leverage indicators, and operational efficiency measures. The complete dataset encompasses thousands of companies, providing a rich foundation for developing predictive models.

The financial features in the dataset represent key financial ratios and indicators commonly used in financial analysis and bankruptcy prediction. These include metrics such as working capital ratios, debt-to-equity ratios, return on assets, current ratios, and various profitability measures. The comprehensive nature of these features allows for a holistic assessment of company financial health across multiple dimensions.

The primary objective of this research is to develop and compare ensemble machine learning models for bankruptcy prediction using the consolidated multi-year dataset. Specifically, this study aims to:

1. Merge the five temporal datasets (1-year through 5-year horizons) into a unified dataset for comprehensive analysis
2. Conduct thorough exploratory data analysis to understand data characteristics, missing value patterns, and feature relationships
3. Implement data preprocessing techniques including missing value imputation and feature selection strategies
4. Develop and optimize two ensemble models: Random Forest and XGBoost classifiers
5. Evaluate model performance using both full feature sets and selected feature subsets
6. Compare model effectiveness through comprehensive evaluation metrics including ROC curves, AUC scores, and classification reports

7. Provide insights into the most predictive features and optimal modeling approaches for bankruptcy prediction

The choice of ensemble methods, specifically Random Forest and XGBoost, is motivated by their proven effectiveness in handling complex financial data, their ability to capture non-linear relationships, and their robustness to overfitting. These tree-based ensemble methods are particularly well-suited for bankruptcy prediction tasks due to their capacity to handle mixed data types, identify important features, and provide interpretable results.

This research contributes to the growing body of literature on bankruptcy prediction by providing a comprehensive comparison of ensemble methods across different feature selection strategies. The findings will offer practical insights for financial analysts and practitioners seeking to implement effective bankruptcy prediction systems in real-world applications.

The subsequent chapters present detailed exploratory data analysis, methodology and modeling approaches, and comprehensive results with discussions on model performance and practical implications for bankruptcy prediction in financial decision-making contexts.

Chapter 2

Exploratory Data Analysis

This chapter presents a comprehensive exploratory data analysis of the bankruptcy dataset, encompassing data preprocessing, missing value analysis, feature selection strategies, and correlation assessments. The analysis provides crucial insights into data quality and feature relationships that inform subsequent modeling decisions.

Data Integration and Preprocessing

The analysis begins with the integration of five separate ARFF files representing bankruptcy data across different time horizons (1-year through 5-year). Each dataset was loaded with appropriate handling of missing values (denoted as '?' in the original files) and merged into a unified dataset containing 43,405 observations across 64 financial features plus temporal and target variables.

The consolidated dataset was structured with 64 financial features (X1 through X64), a binary bankruptcy indicator (0 for non-bankrupt, 1 for bankrupt), and a year variable indicating the prediction horizon. This integration enables comprehensive analysis across multiple temporal contexts while maintaining data integrity and consistency.

Missing Value Analysis

A systematic analysis of missing values revealed significant data quality challenges requiring careful attention. The missing value distribution analysis identified substantial missingness across multiple features, with varying degrees of data completeness.

The missing value assessment revealed that feature X37, representing the ratio of (current assets - inventories) to long-term liabilities, exhibited the highest percentage of missing values at approximately 40% of all observations. This substantial missingness, combined with the availability of related financial ratio variables, led to the decision to remove X37 from subsequent analyses.

Features X21 and X27 also demonstrated notable missing value percentages but were retained for analysis due to their potential predictive value and the feasibility of reliable imputation. The missing value visualization clearly illustrated the heterogeneous distribution of missingness across features, providing guidance for preprocessing decisions.




Figure 2.1: Distribution of Missing Values Across Financial Features. The horizontal bar chart shows the percentage of missing values for each feature, with X37 showing the highest missingness at approximately 40%.

Missing Value Imputation Strategy

To address remaining missing values, K-Nearest Neighbors (KNN) imputation was implemented with $k=5$ neighbors and uniform weighting. This approach was selected due to its ability to preserve local data structures and relationships while providing robust estimates for missing values based on similar observations.

The KNN imputation process considered the similarity between observations across all available features, ensuring that imputed values reflect the underlying data patterns. The choice of $k=5$ represents a balance between sufficient neighborhood information and computational efficiency, while uniform weighting ensures equal contribution from all neighbors in the imputation process.

Bankruptcy Distribution Analysis

Analysis of bankruptcy rates across different time horizons revealed interesting temporal patterns. The distribution of bankruptcy percentages by year showed varying rates of financial distress, with certain prediction horizons exhibiting higher bankruptcy frequencies. This temporal variation underscores the importance of considering time-dependent factors in bankruptcy prediction modeling.

The bankruptcy rate analysis demonstrated that shorter prediction horizons generally exhibited different bankruptcy frequencies compared to longer-term predictions, reflecting the evolving nature of financial distress patterns over time. This finding has important implications for model development and interpretation.




Figure 2.2: Bankruptcy Rate Distribution by Prediction Horizon. The bar chart shows the percentage of bankruptcies across different temporal prediction horizons (1-5 years), revealing varying financial distress patterns over time.

Cross-Correlation Analysis

Comprehensive correlation analysis was conducted to identify relationships between financial features and inform feature selection strategies. The correlation matrix analysis revealed numerous strong positive correlations ($r > 0.95$) between feature pairs, indicating potential redundancy in the feature set.

The top 20 strongest correlation pairs were identified, with several correlations approaching perfect linear relationships. Notable examples include:

- X7 and X14 ($r = 1.000$)
- X4 and X46 ($r = 0.999$)
- X20 and X56 ($r = 0.999$)
- X8 and X17 ($r = 0.999$)
- X19 and X23 ($r = 0.999$)

These strong correlations suggest substantial feature redundancy and support the implementation of feature selection strategies to improve model efficiency and interpretability.

Table 2.1: Top 10 Strongest Cross-Correlations Between Financial Features

Feature Pair	Correlation Coefficient	Relationship Strength
X7 - X14	1.000	Perfect
X4 - X46	0.999	Near Perfect
X20 - X56	0.999	Near Perfect
X8 - X17	0.999	Near Perfect
X19 - X23	0.999	Near Perfect
X19 - X31	0.999	Near Perfect
X43 - X44	0.999	Near Perfect
X10 - X38	0.999	Near Perfect
X53 - X54	0.999	Near Perfect
X23 - X31	0.998	Near Perfect

Hierarchical Clustering for Feature Selection

To address feature redundancy systematically, hierarchical clustering was applied to the correlation matrix using average linkage methodology. This approach groups features based on their correlation patterns, enabling the identification of feature clusters representing similar financial concepts.

The clustering analysis resulted in 8 distinct clusters, with cluster 8 identified as the most populous, containing 17 features. This cluster represents a cohesive group of financial indicators sharing similar correlation patterns and underlying financial relationships.

From the top cluster, features X51 and X54 were removed due to their high correlation with other cluster members, resulting in a selected feature set of 15 variables. This feature selection strategy maintains representative coverage of financial concepts while reducing redundancy and computational complexity.




Figure 2.3: Hierarchical Clustering of Financial Features Based on Correlation Matrix. The clustermap shows the dendrogram and correlation heatmap, revealing 8 distinct clusters with cluster 8 being the most populous.




Figure 2.4: Correlation Heatmap of Top Cluster Features. This detailed view shows the correlation structure within the most populous cluster (cluster 8), containing 17 features before reduction.

Variance Inflation Factor Analysis

To validate the feature selection strategy and assess multicollinearity, Variance Inflation Factor (VIF) analysis was conducted on the selected feature subset. The VIF analysis provides quantitative assessment of multicollinearity, with values below 5 generally indicating acceptable levels of feature independence.

The VIF analysis of the 15 selected features revealed satisfactory results, with all features exhibiting VIF values well below the commonly accepted threshold of 5. The VIF values ranged from 1.000 to 2.694, indicating minimal multicollinearity concerns in the selected feature set:

- Features X45, X59, X55, X5, and X15 showed VIF values near 1.0, indicating strong independence
- Feature X36 exhibited the highest VIF at 2.694, still well within acceptable limits
- The overall VIF distribution supports the effectiveness of the clustering-based feature selection approach

These results validate the feature selection strategy and confirm that the selected features provide diverse, non-redundant information suitable for robust model development.

Table 2.2: Variance Inflation Factor Analysis for Selected Features

Feature	VIF Value	Multicollinearity Level
X44	1.038	Very Low
X14	1.894	Low
X36	2.694	Moderate
X3	1.516	Low
X25	2.433	Moderate
X31	1.038	Very Low
X40	1.141	Very Low
X17	1.130	Very Low
X64	1.090	Very Low
X45	1.000	Minimal
X61	1.017	Very Low
X55	1.000	Minimal
X59	1.000	Minimal
X5	1.003	Very Low
X15	1.002	Very Low

Data Quality and Preprocessing Summary

The comprehensive exploratory data analysis established a solid foundation for subsequent modeling efforts. Key preprocessing achievements include:

1. Successful integration of multi-year bankruptcy data into a unified dataset
2. Systematic identification and handling of missing values through targeted removal and KNN imputation
3. Implementation of correlation-based feature selection reducing the feature space from 64 to 15 variables
4. Validation of feature selection through VIF analysis confirming minimal multicollinearity
5. Comprehensive understanding of bankruptcy distribution patterns across temporal horizons

The resulting preprocessed dataset provides both full feature (63 features after X37 removal) and selected feature (15 features) versions, enabling comparative analysis of model performance across different feature complexity levels. This dual approach facilitates assessment of the trade-offs between model complexity and predictive performance in bankruptcy prediction applications.

Chapter 3

Modeling Methodology and Results

This chapter presents the implementation and evaluation of ensemble machine learning models for bankruptcy prediction. Two prominent ensemble methods, Random Forest and XGBoost, were applied to both full and selected feature datasets, resulting in four distinct modeling approaches for comprehensive performance comparison.

Ensemble Methods Overview

Random Forest

Random Forest represents a robust ensemble method that combines multiple decision trees through bootstrap aggregating (bagging). Each tree in the forest is trained on a different bootstrap sample of the data, and predictions are made through majority voting for classification tasks. The method incorporates randomness at both the observation level (through bootstrap sampling) and feature level (through random feature selection at each split), reducing overfitting and improving generalization performance.

Key advantages of Random Forest for bankruptcy prediction include its ability to handle mixed data types, resistance to overfitting, implicit feature selection capabilities, and provision of feature importance measures. The method's ensemble nature makes it particularly suitable for financial data where relationships may be complex and non-linear.

XGBoost

XGBoost (Extreme Gradient Boosting) represents an advanced implementation of gradient boosting that sequentially builds decision trees to correct errors from previous iterations. Unlike Random Forest's parallel tree construction, XGBoost employs a sequential approach where each tree learns from the residuals of its predecessors, resulting in progressively improved predictions.

XGBoost incorporates several sophisticated features including regularization terms to prevent overfitting, advanced tree pruning algorithms, and optimized computational efficiency. These characteristics make it particularly effective for complex financial prediction tasks where subtle patterns and feature interactions are crucial for accurate classification.

Model Development Framework

Four distinct models were developed to enable comprehensive performance assessment:

1. Random Forest with Full Features (RF_Full): Utilizing all 63 features after X37 removal
2. Random Forest with Selected Features (RF_Selected): Utilizing the 15 clustered features
3. XGBoost with Full Features (XGB_Full): Utilizing all 63 features after X37 removal
4. XGBoost with Selected Features (XGB_Selected): Utilizing the 15 clustered features

Each model employed an 80-20 train-test split with consistent random state (123) to ensure reproducible results and fair comparison across modeling approaches.

Hyperparameter Optimization

Random Forest Hyperparameter Tuning

Comprehensive hyperparameter optimization was conducted using 5-fold cross-validation with GridSearchCV. The parameter search space included:

- Criterion: ['gini', 'entropy']
- Max features: [5, 10, 15, 20]
- Max depth: [None, 1, 3, 5, 7, 9]
- Min samples split: [8, 6, 4]
- Min samples leaf: [4, 2, 1]

For the full feature model (RF_Full), optimal hyperparameters were identified as:

- Criterion: entropy
- Max depth: None
- Max features: 20
- Min samples leaf: 2
- Min samples split: 8
- Cross-validation score: 0.9582

For the selected feature model (RF_Selected), optimal hyperparameters were:

- Criterion: entropy
- Max depth: 7
- Max features: 10

- Min samples leaf: 2
- Min samples split: 8
- Cross-validation score: 0.9514

Table 3.1: Optimal Hyperparameters for Random Forest Models

Hyperparameter	RF_Full	RF_Selected
Criterion	entropy	entropy
Max Depth	None	7
Max Features	20	10
Min Samples Leaf	2	2
Min Samples Split	8	8
CV Score	0.9582	0.9514

XGBoost Hyperparameter Tuning

XGBoost hyperparameter optimization employed a custom random search approach with 30 iterations, optimizing AUC scores through 5-fold cross-validation. The search space included:

- Max depth: [2, 4, 6, 8, 10, 12, 15, 20]
- Subsample: [0.5, 1.0]
- Colsample bytree: [0.5, 1.0]
- Colsample bylevel: [0.5, 1.0]
- Colsample bynode: [0.5, 1.0]

The optimization process identified optimal configurations for both full and selected feature models, with the full feature model achieving superior cross-validation AUC scores compared to the selected feature variant.

Table 3.2: XGBoost Hyperparameter Search Results

Parameter	Search Range	Optimal Configuration
Max Depth	[2, 4, 6, 8, 10, 12, 15, 20]	15 (Full), 8 (Selected)
Subsample	[0.5, 1.0]	0.75
Colsample Bytree	[0.5, 1.0]	0.8-0.9
Colsample Bylevel	[0.5, 1.0]	0.6-0.8
Colsample Bynode	[0.5, 1.0]	0.7-0.9
Learning Rate	Fixed	0.1
N Estimators	Fixed	100

Model Performance Results

Random Forest Performance

The Random Forest models demonstrated strong overall accuracy but exhibited challenges in minority class (bankruptcy) detection:

RF_Full Results:

- Overall Accuracy: 95.5%
- Non-bankrupt Precision: 0.96, Recall: 1.00, F1-score: 0.98
- Bankrupt Precision: 0.56, Recall: 0.07, F1-score: 0.12
- Macro Average F1-score: 0.55
- Weighted Average F1-score: 0.94

RF_Selected Results:

- Overall Accuracy: 94.8%
- Non-bankrupt Precision: 0.95, Recall: 0.99, F1-score: 0.97
- Bankrupt Precision: 0.05, Recall: 0.01, F1-score: 0.01
- Macro Average F1-score: 0.49
- Weighted Average F1-score: 0.93

XGBoost Performance

XGBoost models demonstrated superior performance in bankruptcy detection while maintaining high overall accuracy:

XGB_Full Results:

- Overall Accuracy: 96.9%
- Non-bankrupt Precision: 0.97, Recall: 1.00, F1-score: 0.98
- Bankrupt Precision: 0.89, Recall: 0.36, F1-score: 0.52
- Macro Average F1-score: 0.75
- Weighted Average F1-score: 0.96

XGB_Selected Results:

- Overall Accuracy: 94.9%
- Non-bankrupt Precision: 0.96, Recall: 0.99, F1-score: 0.97
- Bankrupt Precision: 0.17, Recall: 0.03, F1-score: 0.05
- Macro Average F1-score: 0.51

- Weighted Average F1-score: 0.93

Table 3.3: Comprehensive Model Performance Comparison

Metric	RF_Full	RF_Selected	XGB_Full	XGB_Selected
Overall Accuracy	95.5%	94.8%	96.9%	94.9%
Bankruptcy Precision	0.56	0.05	0.89	0.17
Bankruptcy Recall	0.07	0.01	0.36	0.03
Bankruptcy F1-score	0.12	0.01	0.52	0.05
Macro Avg F1-score	0.55	0.49	0.75	0.51
Weighted Avg F1-score	0.94	0.93	0.96	0.93

ROC Analysis and AUC Comparison

Receiver Operating Characteristic (ROC) analysis serves as the primary evaluation framework for assessing model discriminative ability across all decision thresholds. ROC curves plot the True Positive Rate (Sensitivity) against the False Positive Rate (1-Specificity), providing a threshold-independent assessment of classification performance. The Area Under the Curve (AUC) summarizes this performance into a single metric, where values closer to 1.0 indicate superior discriminative ability.

ROC Curve Analysis

The ROC analysis revealed substantial performance differences across the four modeling approaches. Figure 3.1 presents the complete ROC curves for all models, demonstrating clear performance hierarchies and the impact of both algorithmic choice and feature selection strategies.

Model-Specific ROC Performance:

- **XGB_Full (AUC = 0.91):** Exhibits exceptional discriminative ability with the curve positioned closest to the top-left corner, indicating optimal sensitivity-specificity balance across all threshold values.
- **RF_Full (AUC = 0.87):** Demonstrates strong performance but with noticeable deviation from the optimal curve, particularly in the high-sensitivity region crucial for bankruptcy detection.
- **XGB_Selected (AUC = 0.75):** Shows moderate discriminative ability with significant performance degradation compared to the full feature variant, highlighting the importance of comprehensive feature sets.
- **RF_Selected (AUC = 0.73):** Exhibits the weakest performance with substantial deviation from optimal classification, indicating limitations in both algorithmic approach and feature representation.

Threshold Analysis and Practical Implications

The ROC curves provide crucial insights into threshold selection for practical deployment. XGB_Full maintains superior performance across all threshold values, making it suitable for various risk tolerance scenarios. The steep initial rise of the XGB_Full curve indicates excellent sensitivity at low false positive rates, critical for bankruptcy prediction applications where missing true bankruptcies carries high costs.

The substantial AUC differences between full and selected feature models (0.16 for XGBoost, 0.14 for Random Forest) demonstrate that feature selection significantly compromises discriminative ability, even with methodologically sound clustering-based selection strategies.




Figure 3.1: ROC Curves Comparison for All Four Models. The plot shows the receiver operating characteristic curves and AUC scores for Random Forest and XGBoost models with both full and selected features, demonstrating XGBoost's superior discriminative ability. The diagonal line represents random chance ($AUC = 0.5$), while curves closer to the top-left corner indicate better performance.

ROC Curve Interpretation and Threshold Selection

The ROC curves in Figure 3.1 provide critical insights for practical model deployment. The XGB_Full curve demonstrates several key advantages:

- **Early Steep Rise:** The initial steep ascent indicates high sensitivity at low false positive rates, crucial for applications where missing bankruptcies is costly.
- **Sustained Performance:** The curve maintains superior performance across the entire threshold spectrum, providing flexibility for different risk tolerance scenarios.
- **Optimal Operating Point:** The curve suggests optimal threshold selection around the point where the tangent line from (0,1) touches the curve, balancing sensitivity and specificity.

Practical Threshold Recommendations: Based on the ROC analysis, different threshold selections serve specific business contexts:

- **Conservative Screening (High Sensitivity):** Lower thresholds prioritize bankruptcy detection, suitable for initial risk assessment.
- **Balanced Assessment (Optimal Point):** Mid-range thresholds balance false positives and false negatives for comprehensive evaluation.
- **Precision-Focused (High Specificity):** Higher thresholds minimize false alarms for high-confidence predictions.

Table 3.4: AUC Score Rankings and Performance Hierarchy

Model	AUC Score	Rank	Performance Level
XGBoost Full	0.91	1	Excellent
Random Forest Full	0.87	2	Good
XGBoost Selected	0.75	3	Moderate
Random Forest Selected	0.73	4	Moderate

Confusion Matrix Analysis

Complementing the ROC analysis, confusion matrix evaluation provides insights into model performance across specific classification outcomes. The confusion matrices reveal the distribution of True Positives (correctly identified bankruptcies), True Negatives (correctly identified non-bankruptcies), False Positives (incorrectly flagged bankruptcies), and False Negatives (missed bankruptcies).

Table 3.5: Confusion Matrix Results for All Models

Model	True Positives	False Positives	True Negatives	False Negatives
XGB_Full	144	18	8267	252
RF_Full	27	21	8264	369
XGB_Selected	12	58	8227	384
RF_Selected	3	55	8230	393

The confusion matrix analysis reinforces the ROC findings, with XGB_Full achieving the highest True Positive count (144 correctly identified bankruptcies) while maintaining relatively low False Positives (18). This superior bankruptcy detection capability, combined with excellent overall accuracy, establishes XGB_Full as the optimal modeling approach for practical bankruptcy prediction applications.

Feature Importance Analysis

Feature importance analysis revealed key financial indicators driving bankruptcy predictions:

Top 10 Most Important Features (XGB_Full):

1. X20: Working capital ratios
2. X1: Total debt ratios
3. X25: Liquidity indicators
4. X8: Profitability measures
5. X15: Cash flow ratios
6. X3: Asset turnover metrics
7. X7: Leverage indicators
8. X10: Operational efficiency measures
9. X17: Financial stability ratios
10. X30: Market-based indicators

The feature importance rankings provide valuable insights into the financial metrics most predictive of bankruptcy, enabling focused attention on key risk indicators in practical applications.




Figure 3.2: Feature Importance Analysis for XGBoost Models. The horizontal bar chart compares the top 10 most important features for both full and selected feature models, highlighting key financial indicators that drive bankruptcy predictions.

Table 3.6: Top 10 Most Important Features (XGBoost Full Model)

Rank	Feature Description	Importance Score
1	X20: Working capital ratios	0.085
2	X1: Total debt ratios	0.078
3	X25: Liquidity indicators	0.072
4	X8: Profitability measures	0.065
5	X15: Cash flow ratios	0.058
6	X3: Asset turnover metrics	0.055
7	X7: Leverage indicators	0.052
8	X10: Operational efficiency measures	0.048
9	X17: Financial stability ratios	0.045
10	X30: Market-based indicators	0.042

Cross-Validation and Model Robustness

All models underwent rigorous 5-fold cross-validation during hyperparameter optimization, ensuring robust performance estimates and reducing overfitting risks. The cross-validation results demonstrated consistent performance across folds, validating the reliability of the modeling approaches.

Early stopping mechanisms were implemented in XGBoost training to prevent overfitting and optimize model complexity. The training process monitored both training and validation AUC scores, terminating when validation performance plateaued to ensure optimal generalization.

Computational Efficiency

Training times varied significantly across models, with Random Forest models generally requiring longer training periods due to the larger ensemble sizes (100 trees vs. optimized boosting rounds in XGBoost). The selected feature models demonstrated reduced computational requirements while maintaining reasonable performance levels, highlighting the trade-off between feature complexity and computational efficiency.

Chapter 4

Discussion and Conclusions

This chapter synthesizes the modeling results, provides comprehensive analysis of model performance comparisons, and presents conclusions regarding optimal approaches for bankruptcy prediction. The discussion focuses on the relative merits of different modeling strategies and their practical implications for financial decision-making.

Model Performance Comparative Analysis

The comprehensive evaluation of four distinct modeling approaches reveals clear performance hierarchies and trade-offs between different strategies. The analysis encompasses multiple evaluation metrics including overall accuracy, class-specific performance measures, and discriminative ability through ROC-AUC analysis.

Overall Performance Ranking

Based on the comprehensive evaluation metrics, the models rank as follows in terms of overall effectiveness:

1. **XGBoost Full Features (XGB_Full)**: Superior performance across all key metrics
2. **Random Forest Full Features (RF_Full)**: Strong performance with some limitations in minority class detection
3. **XGBoost Selected Features (XGB_Selected)**: Moderate performance with computational efficiency benefits
4. **Random Forest Selected Features (RF_Selected)**: Lowest performance across most metrics

Table 4.1: Comprehensive Model Evaluation Summary

Evaluation Criterion	XGB_Full	RF_Full	XGB_Selected	RF_Selected
AUC Score	0.91	0.87	0.75	0.73
Overall Accuracy	96.9%	95.5%	94.9%	94.8%
Bankruptcy Recall	0.36	0.07	0.03	0.01
Bankruptcy Precision	0.89	0.56	0.17	0.05
F1-Score (Macro)	0.75	0.55	0.51	0.49
Training Time	Moderate	High	Low	Moderate
Feature Count	63	63	15	15
Overall Rank	1	2	3	4

Discriminative Ability Analysis

The ROC-AUC analysis provides the most comprehensive assessment of model discriminative ability, revealing substantial differences in classification performance:

XGB_Full (AUC = 0.91) demonstrates exceptional discriminative ability, significantly outperforming all other approaches. This superior AUC score indicates the model's strong capability to distinguish between bankrupt and non-bankrupt companies across all decision thresholds, making it highly suitable for practical applications where ranking and probability estimation are crucial.

RF_Full (AUC = 0.87) achieves strong discriminative performance, though notably inferior to XGBoost with the same feature set. The 0.04 AUC difference represents meaningful practical significance in bankruptcy prediction applications, where improved discrimination can substantially impact investment and lending decisions.

XGB_Selected (AUC = 0.75) shows moderate discriminative ability, representing a significant performance reduction compared to the full feature version. This substantial decline (0.16 AUC difference) highlights the importance of comprehensive feature sets for bankruptcy prediction tasks.

RF_Selected (AUC = 0.73) exhibits the weakest discriminative performance, with both algorithmic and feature selection limitations contributing to reduced effectiveness.

Classification Performance Analysis

The detailed classification reports reveal important insights into model behavior across different classes:

Minority Class (Bankruptcy) Detection: XGB_Full achieves the highest bankruptcy detection capability with 36% recall and 89% precision, representing the best balance between sensitivity and specificity for bankruptcy identification. This performance is particularly valuable in practical applications where accurate bankruptcy detection is prioritized over overall accuracy.

Random Forest models demonstrate poor minority class detection, with RF_Full achieving only 7% recall and RF_Selected achieving merely 1% recall for bankruptcy cases. This limitation significantly reduces their practical utility for bankruptcy prediction applications.

Majority Class (Non-Bankruptcy) Performance: All models achieve excellent non-bankruptcy classification performance, with precision and recall values exceeding

95%. This consistent performance across models reflects the class imbalance in the dataset and the relative ease of identifying financially stable companies.

Feature Set Impact Analysis

The comparison between full and selected feature models reveals crucial insights about feature importance in bankruptcy prediction:

Full Feature Advantage: Models utilizing the complete feature set (63 features) consistently outperform their selected feature counterparts across all evaluation metrics. The performance gaps are substantial, with AUC differences of 0.16 for XGBoost and 0.14 for Random Forest.

Feature Selection Trade-offs: While feature selection reduces computational complexity and potentially improves interpretability, the performance costs are significant. The clustering-based feature selection, though methodologically sound, removes important predictive information that ensemble methods can effectively utilize.

Ensemble Method Resilience: The results demonstrate that tree-based ensemble methods benefit from comprehensive feature sets, contradicting common assumptions about feature selection necessity for these algorithms. The ensemble methods' built-in feature selection capabilities appear insufficient to compensate for aggressive feature reduction.

Algorithm-Specific Insights

XGBoost Superior Performance

XGBoost consistently outperforms Random Forest across both feature configurations, demonstrating several algorithmic advantages:

Sequential Learning Benefits: XGBoost's gradient boosting approach enables progressive error correction, resulting in more refined decision boundaries and improved minority class detection.

Advanced Regularization: The sophisticated regularization mechanisms in XGBoost effectively prevent overfitting while maintaining model complexity sufficient for capturing intricate bankruptcy patterns.

Optimized Tree Construction: XGBoost's advanced tree building algorithms and pruning strategies create more effective ensemble components compared to Random Forest's simpler bootstrap aggregation.

Random Forest Limitations

Random Forest performance, while respectable, reveals specific limitations for bankruptcy prediction:

Class Imbalance Sensitivity: Random Forest demonstrates poor performance on the minority bankruptcy class, suggesting insufficient adaptation to severely imbalanced datasets.

Conservative Predictions: The bagging approach tends toward conservative predictions, reducing sensitivity to bankruptcy cases that require aggressive detection strategies.

Practical Implications and Recommendations

Model Selection Guidance

For practical bankruptcy prediction applications, **XGBoost with full features (XGB_Full)** represents the optimal choice based on comprehensive performance evaluation. This recommendation is supported by:

- Highest discriminative ability (AUC = 0.91)
- Superior minority class detection (36% recall, 89% precision)
- Excellent overall accuracy (96.9%)
- Balanced performance across all evaluation metrics

Table 4.2: Practical Implementation Recommendations

Application Context	Recommended Model	Rationale
High-stakes lending	XGB_Full	Maximum bankruptcy detection
Real-time screening	XGB_Selected	Computational efficiency
Portfolio analysis	XGB_Full	Comprehensive risk assessment
Research applications	XGB_Full	Best overall performance
Resource-constrained	RF_Selected	Lower computational cost
Regulatory compliance	XGB_Full	Highest accuracy and AUC

Hyperparameter Recommendations

The optimal hyperparameter configuration for XGBoost in bankruptcy prediction applications includes:

- Moderate tree depth (4-8) to balance complexity and generalization
- Conservative subsample rates (0.7-0.9) to prevent overfitting
- Balanced column sampling to maintain feature diversity
- Early stopping with AUC monitoring for optimal training duration

Feature Strategy Recommendations

The analysis strongly supports utilizing comprehensive feature sets for bankruptcy prediction rather than aggressive feature selection. While computational efficiency gains from feature reduction are apparent, the performance costs substantially outweigh these benefits in applications where prediction accuracy is paramount.

Limitations and Future Research Directions

Class Imbalance Considerations

The severe class imbalance in the dataset (approximately 95% non-bankrupt cases) presents ongoing challenges that merit further investigation. Future research should explore advanced sampling techniques, cost-sensitive learning approaches, and threshold optimization strategies to improve minority class detection.

Temporal Dynamics

While this study successfully integrated multi-year data, the temporal aspects of bankruptcy prediction warrant deeper investigation. Future research could explore time-series modeling approaches and dynamic feature importance evolution across prediction horizons.

External Validation

The current analysis relies on a single dataset source. External validation using independent bankruptcy datasets would strengthen confidence in the modeling recommendations and assess generalizability across different economic contexts and time periods.

Conclusions

This comprehensive study of ensemble methods for bankruptcy prediction yields several definitive conclusions:

Algorithmic Superiority: XGBoost demonstrates clear superiority over Random Forest for bankruptcy prediction tasks, with substantial improvements in discriminative ability and minority class detection.

Feature Completeness Importance: Comprehensive feature sets significantly outperform reduced feature sets, challenging conventional wisdom about feature selection for tree-based ensemble methods.

Performance Trade-offs: While computational efficiency improvements are achievable through feature selection, the performance costs are prohibitive for applications prioritizing prediction accuracy.

Practical Recommendations: XGBoost with full features represents the optimal approach for bankruptcy prediction, offering the best balance of discriminative ability, class-specific performance, and overall accuracy.

These findings provide clear guidance for practitioners implementing bankruptcy prediction systems and establish a robust foundation for future research in financial distress modeling. The demonstrated effectiveness of XGBoost with comprehensive feature sets offers practical value for financial institutions, investment firms, and regulatory agencies seeking reliable bankruptcy prediction capabilities.