

Internet Log2 Data Analysis: A Comprehensive Study of Network Traffic Classification Using Machine Learning Approaches

Advanced Analytics Report

July 7, 2025

Contents

1	Introduction	3
1.1	Dataset Overview	3
1.2	Research Objectives	4
2	Exploratory Data Analysis	4
2.1	Data Quality Assessment	4
2.1.1	Missing Value Analysis	4
2.2	Target Variable Distribution	4
2.3	Categorical Variable Analysis	5
2.3.1	Port Distribution Patterns	5
2.3.2	Top-N Encoding Strategy	5
2.4	Numerical Variable Analysis	6
2.4.1	Correlation Analysis	6
2.4.2	Multicollinearity Assessment using Variance Inflation Factor (VIF)	6
2.5	Feature Engineering and Preprocessing Pipeline	7
3	Machine Learning Modeling	7
3.1	Model Selection Rationale	8
3.2	Support Vector Machine (SVM) Analysis	8
3.2.1	Hyperparameter Optimization	8
3.2.2	Computational Considerations	9
3.3	Stochastic Gradient Descent (SGD) Analysis	9
3.3.1	Model Configuration	9
3.3.2	Cross-Validation Performance	9
3.4	Model Performance Comparison	10
3.4.1	Classification Reports	10
3.4.2	Performance Analysis	10
3.4.3	ROC Curve Analysis	10
3.5	Computational Efficiency Analysis	11
3.5.1	Training Time Comparison	11
3.5.2	Scalability Implications	11

4	Discussion and Conclusions	12
4.1	Key Findings Summary	12
4.1.1	Data Quality and Preprocessing Effectiveness	12
4.1.2	Model Performance Achievement	12
4.1.3	Computational Efficiency Trade-offs	12
4.2	Methodological Contributions	13
4.2.1	Top-N Encoding Validation	13
4.2.2	Linear Model Effectiveness	13
4.3	Practical Implications for Network Security	13
4.3.1	Deployment Recommendations	13
4.3.2	Security Policy Insights	14
4.4	Limitations and Future Research	14
4.4.1	Identified Limitations	14
4.4.2	Future Research Directions	14
4.5	Final Conclusions	15

1 Introduction

Network traffic analysis represents a critical domain in cybersecurity and network management, where understanding and classifying traffic patterns enables organizations to implement effective security policies and optimize network performance. This report presents a comprehensive analysis of Internet Log2 data, focusing on the application of machine learning techniques for traffic classification.

1.1 Dataset Overview

The Internet Log2 dataset, located at `data/log2.csv`, comprises 65,532 network traffic records collected from network monitoring systems. Each record represents a network connection session with comprehensive metadata describing the communication characteristics between source and destination endpoints.

The dataset structure encompasses 12 variables: 11 input features and one target variable. The input features are categorized into two distinct groups:

Categorical Features (4 variables):

- **Source Port:** The originating port number of the network connection
- **Destination Port:** The target port number receiving the connection
- **NAT Source Port:** The Network Address Translation (NAT) source port
- **NAT Destination Port:** The NAT destination port

Numerical Features (7 variables):

- **Bytes:** Total data volume transferred in the connection
- **Bytes Sent:** Volume of data transmitted from source to destination
- **Bytes Received:** Volume of data received from destination to source
- **Packets:** Total number of packets exchanged
- **Elapsed Time (sec):** Duration of the connection session
- **pkts_sent:** Number of packets transmitted from source
- **pkts_received:** Number of packets received at source

Target Variable: The **Action** variable represents the network security decision applied to each connection, categorized into four distinct classes:

- **Allow:** Connections permitted through the network
- **Drop:** Connections silently discarded without notification
- **Deny:** Connections explicitly rejected with notification
- **Reset:** Connections terminated with reset signals

1.2 Research Objectives

This analysis aims to achieve several key objectives:

1. Conduct comprehensive exploratory data analysis to understand traffic patterns and data quality
2. Develop effective preprocessing strategies for handling high-cardinality categorical variables
3. Implement and compare machine learning models for network traffic classification
4. Evaluate model performance and computational efficiency for large-scale deployment
5. Provide actionable insights for network security policy implementation

The significance of this research lies in its practical applications for network security operations, where accurate and efficient traffic classification directly impacts organizational cybersecurity posture and network performance optimization.

2 Exploratory Data Analysis

The exploratory data analysis phase provides crucial insights into data quality, feature characteristics, and underlying patterns that inform subsequent modeling decisions. This comprehensive analysis encompasses missing value assessment, categorical variable examination, numerical variable correlation analysis, and feature engineering considerations.

2.1 Data Quality Assessment

2.1.1 Missing Value Analysis

A critical first step in any data science project involves assessing data completeness. The Internet Log2 dataset demonstrates exceptional data quality, with zero missing values across all 12 variables. This complete dataset eliminates the need for imputation strategies and ensures that all 65,532 records contribute meaningfully to the analysis.

The absence of missing data is particularly valuable in network traffic analysis, where incomplete records could indicate data collection issues or network monitoring gaps that might bias model performance.

2.2 Target Variable Distribution

The distribution of the Action variable reveals significant class imbalance, a common characteristic in network security datasets:

- **Allow:** Approximately 60% of all traffic (dominant class)
- **Drop:** Substantial portion of blocked traffic
- **Deny:** Moderate portion of explicitly rejected connections
- **Reset:** Minimal representation, constituting a very small percentage

This imbalanced distribution reflects real-world network traffic patterns, where legitimate traffic typically dominates over security-blocked connections. The extreme under-representation of the "Reset" class presents particular challenges for model training and evaluation.

2.3 Categorical Variable Analysis

The four port-based categorical variables exhibit dramatically different cardinality characteristics, necessitating sophisticated encoding strategies:

Feature	Unique Values	Total Records	Percentage Unique
Source Port	22,724	65,532	34.68%
Destination Port	3,273	65,532	4.99%
NAT Source Port	29,152	65,532	44.49%
NAT Destination Port	2,533	65,532	3.87%

2.3.1 Port Distribution Patterns

Detailed analysis of the top 10 most frequent ports reveals important networking patterns:

Source and NAT Source Ports:

- Common high-numbered ports (58638, ephemeral range)
- Significant presence of port 0 (placeholder/invalid port indicator)
- Wide distribution suggesting dynamic port allocation

Destination and NAT Destination Ports:

- Port 53: DNS traffic (highly prevalent)
- Port 443: HTTPS/TLS encrypted web traffic
- Port 80: HTTP web traffic
- These standard service ports dominate destination traffic

The prevalence of port 0, particularly in NAT Source Port data, indicates the presence of placeholder or invalid port assignments in the dataset. While port 0 is technically invalid (valid port range: 1-65535), its retention provides valuable information about network configuration patterns.

2.3.2 Top-N Encoding Strategy

Given the extreme cardinality of Source Port (22,724 unique values) and NAT Source Port (29,152 unique values), traditional one-hot encoding would generate prohibitively sparse feature matrices with tens of thousands of columns. This approach would result in:

- Exponential memory requirements
- Computational inefficiency in model training

- Curse of dimensionality problems
- Model interpretability challenges

To address these challenges, we implement a **Top-N encoding strategy** with the following methodology:

1. **Frequency Analysis:** Calculate value frequencies for each categorical variable
2. **Threshold Selection:** Retain the top 1% most frequent values for each variable
3. **Category Consolidation:** Group all remaining values into an "Other" category
4. **One-Hot Encoding:** Apply standard OHE to the reduced categorical sets

This approach balances information preservation with computational efficiency. The 1% threshold was selected through careful consideration of:

- Information retention vs. dimensionality reduction trade-off
- Computational constraints for model training
- Interpretability requirements for business stakeholders

The Top-N encoding reduces the feature space from potentially 57,682 binary features (sum of all unique categorical values) to approximately 580 features (1% of each categorical variable plus "Other" categories), representing a 99% dimensionality reduction while preserving the most informative categorical patterns.

2.4 Numerical Variable Analysis

2.4.1 Correlation Analysis

The correlation matrix for numerical variables reveals critical insights into feature relationships:

- **Strong Positive Correlations:** High correlation between Bytes groups (Bytes, Bytes Sent, Bytes Received) and Packets groups (Packets, pkts_sent, pkts_received)
- **Architectural Relationships:** Bytes = Bytes Sent + Bytes Received; Packets = pkts_sent + pkts_received
- **Independence:** Elapsed Time demonstrates relative independence from other variables

2.4.2 Multicollinearity Assessment using Variance Inflation Factor (VIF)

Variance Inflation Factor analysis quantifies the severity of multicollinearity among numerical predictors:

Initial VIF Analysis (All Variables):

- Bytes, Bytes Sent, Bytes Received: VIF = infinity (perfect multicollinearity)
- Packets, pkts_sent, pkts_received: VIF = infinity (perfect multicollinearity)

- Elapsed Time: $VIF = 1.02$ (acceptable)

The infinite VIF values confirm perfect linear relationships among byte and packet variables, necessitating feature selection to avoid numerical instability in model training.

Refined VIF Analysis (Selected Variables):

After retaining only aggregated variables (Bytes, Packets, Elapsed Time):

- Bytes: $VIF = 19.78$
- Packets: $VIF = 19.78$
- Elapsed Time: $VIF = 1.02$

The reduced VIF values of approximately 20 for Bytes and Packets, while elevated, remain within acceptable ranges for model training. These variables capture the essential information content while eliminating perfect multicollinearity.

2.5 Feature Engineering and Preprocessing Pipeline

The final preprocessing pipeline integrates categorical and numerical transformations:

1. Categorical Processing:

- Apply Top-N encoding (1% threshold) to all port variables
- Generate "Other" categories for low-frequency values
- Apply one-hot encoding to transformed categorical variables

2. Numerical Processing:

- Select Bytes, Packets, and Elapsed Time variables
- Apply StandardScaler for feature normalization
- Ensure zero mean and unit variance for numerical features

3. Feature Integration:

- Concatenate processed categorical and numerical features
- Verify feature matrix dimensions and data types
- Prepare final dataset for model training

This comprehensive preprocessing approach results in a final feature matrix with approximately 580 features, representing a substantial dimensionality reduction from the original high-cardinality categorical space while preserving essential information content for accurate traffic classification.

3 Machine Learning Modeling

This section presents a comprehensive comparison of Support Vector Machine (SVM) and Stochastic Gradient Descent (SGD) classifiers for network traffic classification. The analysis encompasses hyperparameter optimization, cross-validation assessment, computational efficiency evaluation, and performance comparison through multiple metrics.

3.1 Model Selection Rationale

The selection of SVM and SGD classifiers is motivated by several factors relevant to network traffic classification:

- **SVM Strengths:** Excellent performance on medium-sized datasets, robust handling of high-dimensional feature spaces, and strong theoretical foundations
- **SGD Advantages:** Computational efficiency for large datasets, scalability to massive data volumes, and comparable performance to more complex algorithms
- **Linear Separability:** Initial analysis suggests that network traffic classes may be linearly separable, making linear models appropriate
- **Interpretability:** Both models provide insights into feature importance for security policy development

3.2 Support Vector Machine (SVM) Analysis

3.2.1 Hyperparameter Optimization

SVM performance was evaluated across multiple regularization parameter (C) values using 5-fold cross-validation. The comprehensive hyperparameter search examined C values spanning four orders of magnitude:

C Parameter	Mean CV Score
0.001	0.9878
0.01	0.9951
0.1	0.9951
1.0	0.9960
10.0	0.9965
100.0	0.9968
1000.0	0.9970

The results demonstrate several key insights:

- **Performance Plateau:** CV scores plateau above $C = 0.01$, suggesting optimal regularization
- **Marginal Improvements:** Higher C values yield diminishing returns in accuracy
- **Stability:** Consistent performance across the optimal range indicates robust model behavior
- **Linear Kernel Effectiveness:** The linear kernel achieves exceptional performance (99.5%), eliminating the need for complex polynomial or RBF kernels

3.2.2 Computational Considerations

SVM training on this dataset (65,532 samples $\times$ 580 features) presents significant computational challenges:

- **Quadratic Complexity:** SVM optimization scales approximately $O(n^2)$ to $O(n^3)$ with sample size
- **Memory Requirements:** Kernel matrix storage demands substantial memory resources
- **Training Time:** Extended training periods for hyperparameter optimization
- **Scalability Limitations:** Performance degradation with larger datasets

Despite achieving excellent accuracy, SVM's computational demands limit its applicability for real-time network traffic classification in high-volume environments.

3.3 Stochastic Gradient Descent (SGD) Analysis

3.3.1 Model Configuration

The SGD classifier was configured with optimization-focused parameters:

- **Loss Function:** Log-loss for probabilistic output compatibility
- **Maximum Iterations:** 1,000 with early stopping criteria
- **Validation Strategy:** 10% validation split for convergence monitoring
- **Convergence Tolerance:** $1e-3$ for stable optimization
- **Early Stopping:** 5 iterations without improvement threshold

3.3.2 Cross-Validation Performance

SGD cross-validation results demonstrate competitive performance with superior computational efficiency:

CV Fold	Accuracy Score
Fold 1	0.9947
Fold 2	0.9947
Fold 3	0.9951
Fold 4	0.9950
Fold 5	0.9952
Mean CV Score	0.9949

Computational Performance: SGD cross-validation completed in 3.09 seconds (wall time), representing a dramatic improvement over SVM training times.

3.4 Model Performance Comparison

3.4.1 Classification Reports

SVM Classification Performance:

Class	Precision	Recall	F1-Score	Support
Allow	1.00	1.00	1.00	7,557
Deny	1.00	0.99	0.99	2,942
Drop	1.00	1.00	1.00	2,601
Reset-both	0.00	0.00	0.00	7
Weighted Avg	1.00	1.00	1.00	13,107

SGD Classification Performance:

Class	Precision	Recall	F1-Score	Support
Allow	1.00	1.00	1.00	7,557
Deny	0.99	1.00	0.99	2,942
Drop	1.00	1.00	1.00	2,601
Reset-both	0.00	0.00	0.00	7
Weighted Avg	1.00	1.00	1.00	13,107

3.4.2 Performance Analysis

Both models demonstrate exceptional performance across major traffic classes:

- **Perfect Classification:** Near-perfect precision and recall for Allow, Drop, and Deny classes
- **Class Imbalance Impact:** Reset-both class shows zero performance metrics due to extreme underrepresentation (7 samples)
- **Comparable Accuracy:** Both models achieve 99.5%+ weighted average performance
- **Minimal Performance Difference:** SGD matches SVM performance with marginal variations

3.4.3 ROC Curve Analysis

Receiver Operating Characteristic (ROC) curves provide additional performance insights through Area Under Curve (AUC) analysis:

SVM ROC-AUC Scores:

- Allow: AUC approximately 1.00
- Deny: AUC approximately 1.00
- Drop: AUC approximately 1.00
- Reset-both: AUC approximately 0.50 (random performance due to sample size)

SGD ROC-AUC Scores:

- Allow: AUC approximately 1.00
- Deny: AUC approximately 1.00
- Drop: AUC approximately 1.00
- Reset-both: AUC approximately 0.50 (random performance due to sample size)

The ROC analysis confirms exceptional discriminative ability for both models across major traffic classes, with identical performance limitations for the severely underrepresented Reset-both class.

3.5 Computational Efficiency Analysis

3.5.1 Training Time Comparison

The computational efficiency comparison reveals dramatic differences:

Operation	SVM Time	SGD Time
5-Fold CV	~30 minutes*	3.09 seconds
Single Model Training	10 minutes*	~1 second
Hyperparameter Search	Hours*	Minutes

*estimated

3.5.2 Scalability Implications

The computational analysis reveals critical scalability considerations:

SVM Limitations:

- Quadratic to cubic scaling with dataset size
- Memory requirements grow quadratically
- Impractical for real-time applications
- Limited applicability to big data scenarios

SGD Advantages:

- Linear scaling with dataset size
- Constant memory requirements
- Online learning capabilities
- Suitable for streaming data applications

These computational characteristics position SGD as the preferred algorithm for large-scale network traffic classification systems, where real-time processing and scalability are paramount concerns.

4 Discussion and Conclusions

This comprehensive analysis of Internet Log2 data demonstrates the successful application of machine learning techniques for network traffic classification, yielding significant insights for both methodological approaches and practical deployment considerations in cybersecurity operations.

4.1 Key Findings Summary

4.1.1 Data Quality and Preprocessing Effectiveness

The Internet Log2 dataset exhibited exceptional data quality with zero missing values across 65,532 records, providing a robust foundation for analysis. The preprocessing pipeline successfully addressed several critical challenges:

- **High-Cardinality Management:** The Top-N encoding strategy effectively reduced dimensionality from 57,682 potential categorical features to approximately 580 features, achieving 99% dimensionality reduction while preserving essential information content.
- **Multicollinearity Resolution:** VIF analysis and feature selection eliminated perfect multicollinearity among numerical variables, reducing infinite VIF values to acceptable ranges (< 20) for model stability.
- **Feature Engineering Success:** The integration of categorical and numerical preprocessing created a balanced feature matrix suitable for machine learning algorithms.

4.1.2 Model Performance Achievement

Both SVM and SGD classifiers demonstrated exceptional performance capabilities:

- **Classification Accuracy:** Both models achieved $>99.5\%$ weighted average accuracy across major traffic classes
- **Class-Specific Performance:** Perfect or near-perfect precision and recall for Allow, Drop, and Deny classes
- **ROC-AUC Excellence:** Area under curve values approaching 1.0 for all major classes
- **Cross-Validation Stability:** Consistent performance across 5-fold cross-validation indicates robust generalization

4.1.3 Computational Efficiency Trade-offs

The analysis revealed dramatic computational differences between algorithms:

- **SVM Limitations:** Training times exceeded 30 minutes for cross-validation, with quadratic scaling constraints limiting scalability

- **SGD Efficiency:** Equivalent performance achieved in 3.09 seconds, representing ~500x computational improvement
- **Scalability Implications:** SGD's linear scaling characteristics enable deployment in high-volume, real-time environments

4.2 Methodological Contributions

4.2.1 Top-N Encoding Validation

This analysis validates the Top-N encoding approach for high-cardinality categorical variables in network traffic analysis:

- **Information Preservation:** The 1% threshold successfully retained the most informative categorical patterns while dramatically reducing feature space complexity
- **Computational Benefits:** Enabled feasible model training on high-cardinality data without sacrificing performance
- **Practical Applicability:** Demonstrates a scalable approach for similar categorical encoding challenges in cybersecurity domains

4.2.2 Linear Model Effectiveness

The success of linear kernels in SVM and linear SGD models provides important insights:

- **Problem Characteristics:** Network traffic classification exhibits linear separability characteristics, eliminating the need for complex non-linear transformations
- **Interpretability Advantages:** Linear models provide transparent decision boundaries facilitating security policy interpretation
- **Computational Efficiency:** Linear models avoid the computational overhead of polynomial or RBF kernel transformations

4.3 Practical Implications for Network Security

4.3.1 Deployment Recommendations

Based on the analysis findings, the following deployment recommendations emerge:

For Real-Time Systems:

- SGD classifiers provide optimal balance of accuracy and computational efficiency
- Linear scaling enables processing of high-volume network traffic streams
- Online learning capabilities support adaptive security policies

For Batch Analysis:

- SVM provides marginal accuracy improvements for offline analysis scenarios
- Suitable for detailed forensic analysis where computational time is less critical
- Beneficial for model validation and benchmark establishment

4.3.2 Security Policy Insights

The model performance across traffic classes provides actionable security insights:

- **High Confidence Classification:** 99% accuracy enables automated policy enforcement for major traffic categories
- **Edge Case Handling:** The Reset-both class requires alternative approaches due to extreme underrepresentation
- **False Positive Minimization:** Excellent precision metrics support deployment in production environments with minimal manual intervention

4.4 Limitations and Future Research

4.4.1 Identified Limitations

Several limitations warrant acknowledgment and future investigation:

- **Class Imbalance:** The Reset-both class's extreme underrepresentation (7 samples) prevents effective classification, requiring specialized sampling or ensemble approaches
- **Top-N Threshold Selection:** The 1% threshold was selected empirically; systematic optimization could yield further improvements
- **Temporal Dynamics:** The analysis lacks temporal components that might reveal evolving attack patterns or network behavior changes
- **Generalization Scope:** Results are specific to this dataset's network environment and may require validation across diverse network infrastructures

4.4.2 Future Research Directions

Several research opportunities emerge from this analysis:

- **Imbalanced Learning:** Investigation of specialized techniques (SMOTE, cost-sensitive learning, ensemble methods) for minority class handling
- **Threshold Optimization:** Systematic optimization of Top-N encoding thresholds through grid search or Bayesian optimization
- **Temporal Analysis:** Incorporation of time-series components to capture dynamic network behavior patterns
- **Ensemble Methods:** Combination of SVM and SGD predictions through ensemble techniques for improved robustness
- **Deep Learning Exploration:** Investigation of neural network approaches for automatic feature learning from raw network data
- **Interpretability Enhancement:** Development of model interpretation techniques for security analyst decision support

4.5 Final Conclusions

This analysis successfully demonstrates the application of machine learning techniques for network traffic classification, achieving several significant outcomes:

1. **Methodology Validation:** The preprocessing pipeline effectively handles real-world network data challenges, providing a replicable framework for similar analyses
2. **Algorithm Comparison:** SGD emerges as the preferred algorithm for practical deployment, offering comparable accuracy with superior computational efficiency
3. **Performance Achievement:** Both models exceed industry-standard accuracy requirements (99%) for automated network security policy enforcement
4. **Scalability Solution:** The Top-N encoding approach provides a scalable solution for high-cardinality categorical data in cybersecurity applications

The research contributes to the growing body of knowledge in applied machine learning for cybersecurity, providing both methodological insights and practical deployment guidance. The exceptional model performance, combined with comprehensive computational analysis, establishes a foundation for operational network traffic classification systems that can enhance organizational cybersecurity capabilities while maintaining computational feasibility for large-scale deployment.

The documented trade-offs between SVM and SGD classifiers provide clear guidance for practitioners selecting algorithms based on their specific operational requirements, whether prioritizing maximum accuracy (SVM) or computational efficiency (SGD) for real-time applications. This analysis ultimately advances the state of practice in machine learning-driven network security operations, demonstrating that sophisticated classification capabilities can be achieved with appropriate preprocessing and algorithm selection strategies.